

IndicCorp v2 corpus provides large monolingual datasets for different Indian languages. The corpus can be accessed from the link below:

<https://huggingface.co/datasets/ai4bharat/IndicCorpV2>

You need to sentence tokenize your data as each row in the corpus is a paragraph, not a sentence.

Download the data of your language and complete the tasks below.

1. You need to create training, development, and test data splits of your data. Your development and test sets should contain 1000 sentences each. Rest all are part of the training data. Create development and test sets by taking length of sentences into account. Your development and test sets should not consist of sentences of similar lengths. [Hint: Pick sentences of different lengths by maintaining a length distribution]
2. You need to implement BPE and WordPiece on your training data and tokenize the development and test data.
3. You need to test on 4 different sizes of training data and show how the tokenization changes with the size.
 - a. 100000
 - b. 300000
 - c. 500000
 - d. 1000000
4. Similarly you need to test on 3 different vocabulary sizes and show the impact of vocabulary on tokenization.
 - a. 20000
 - b. 30000
 - c. 50000